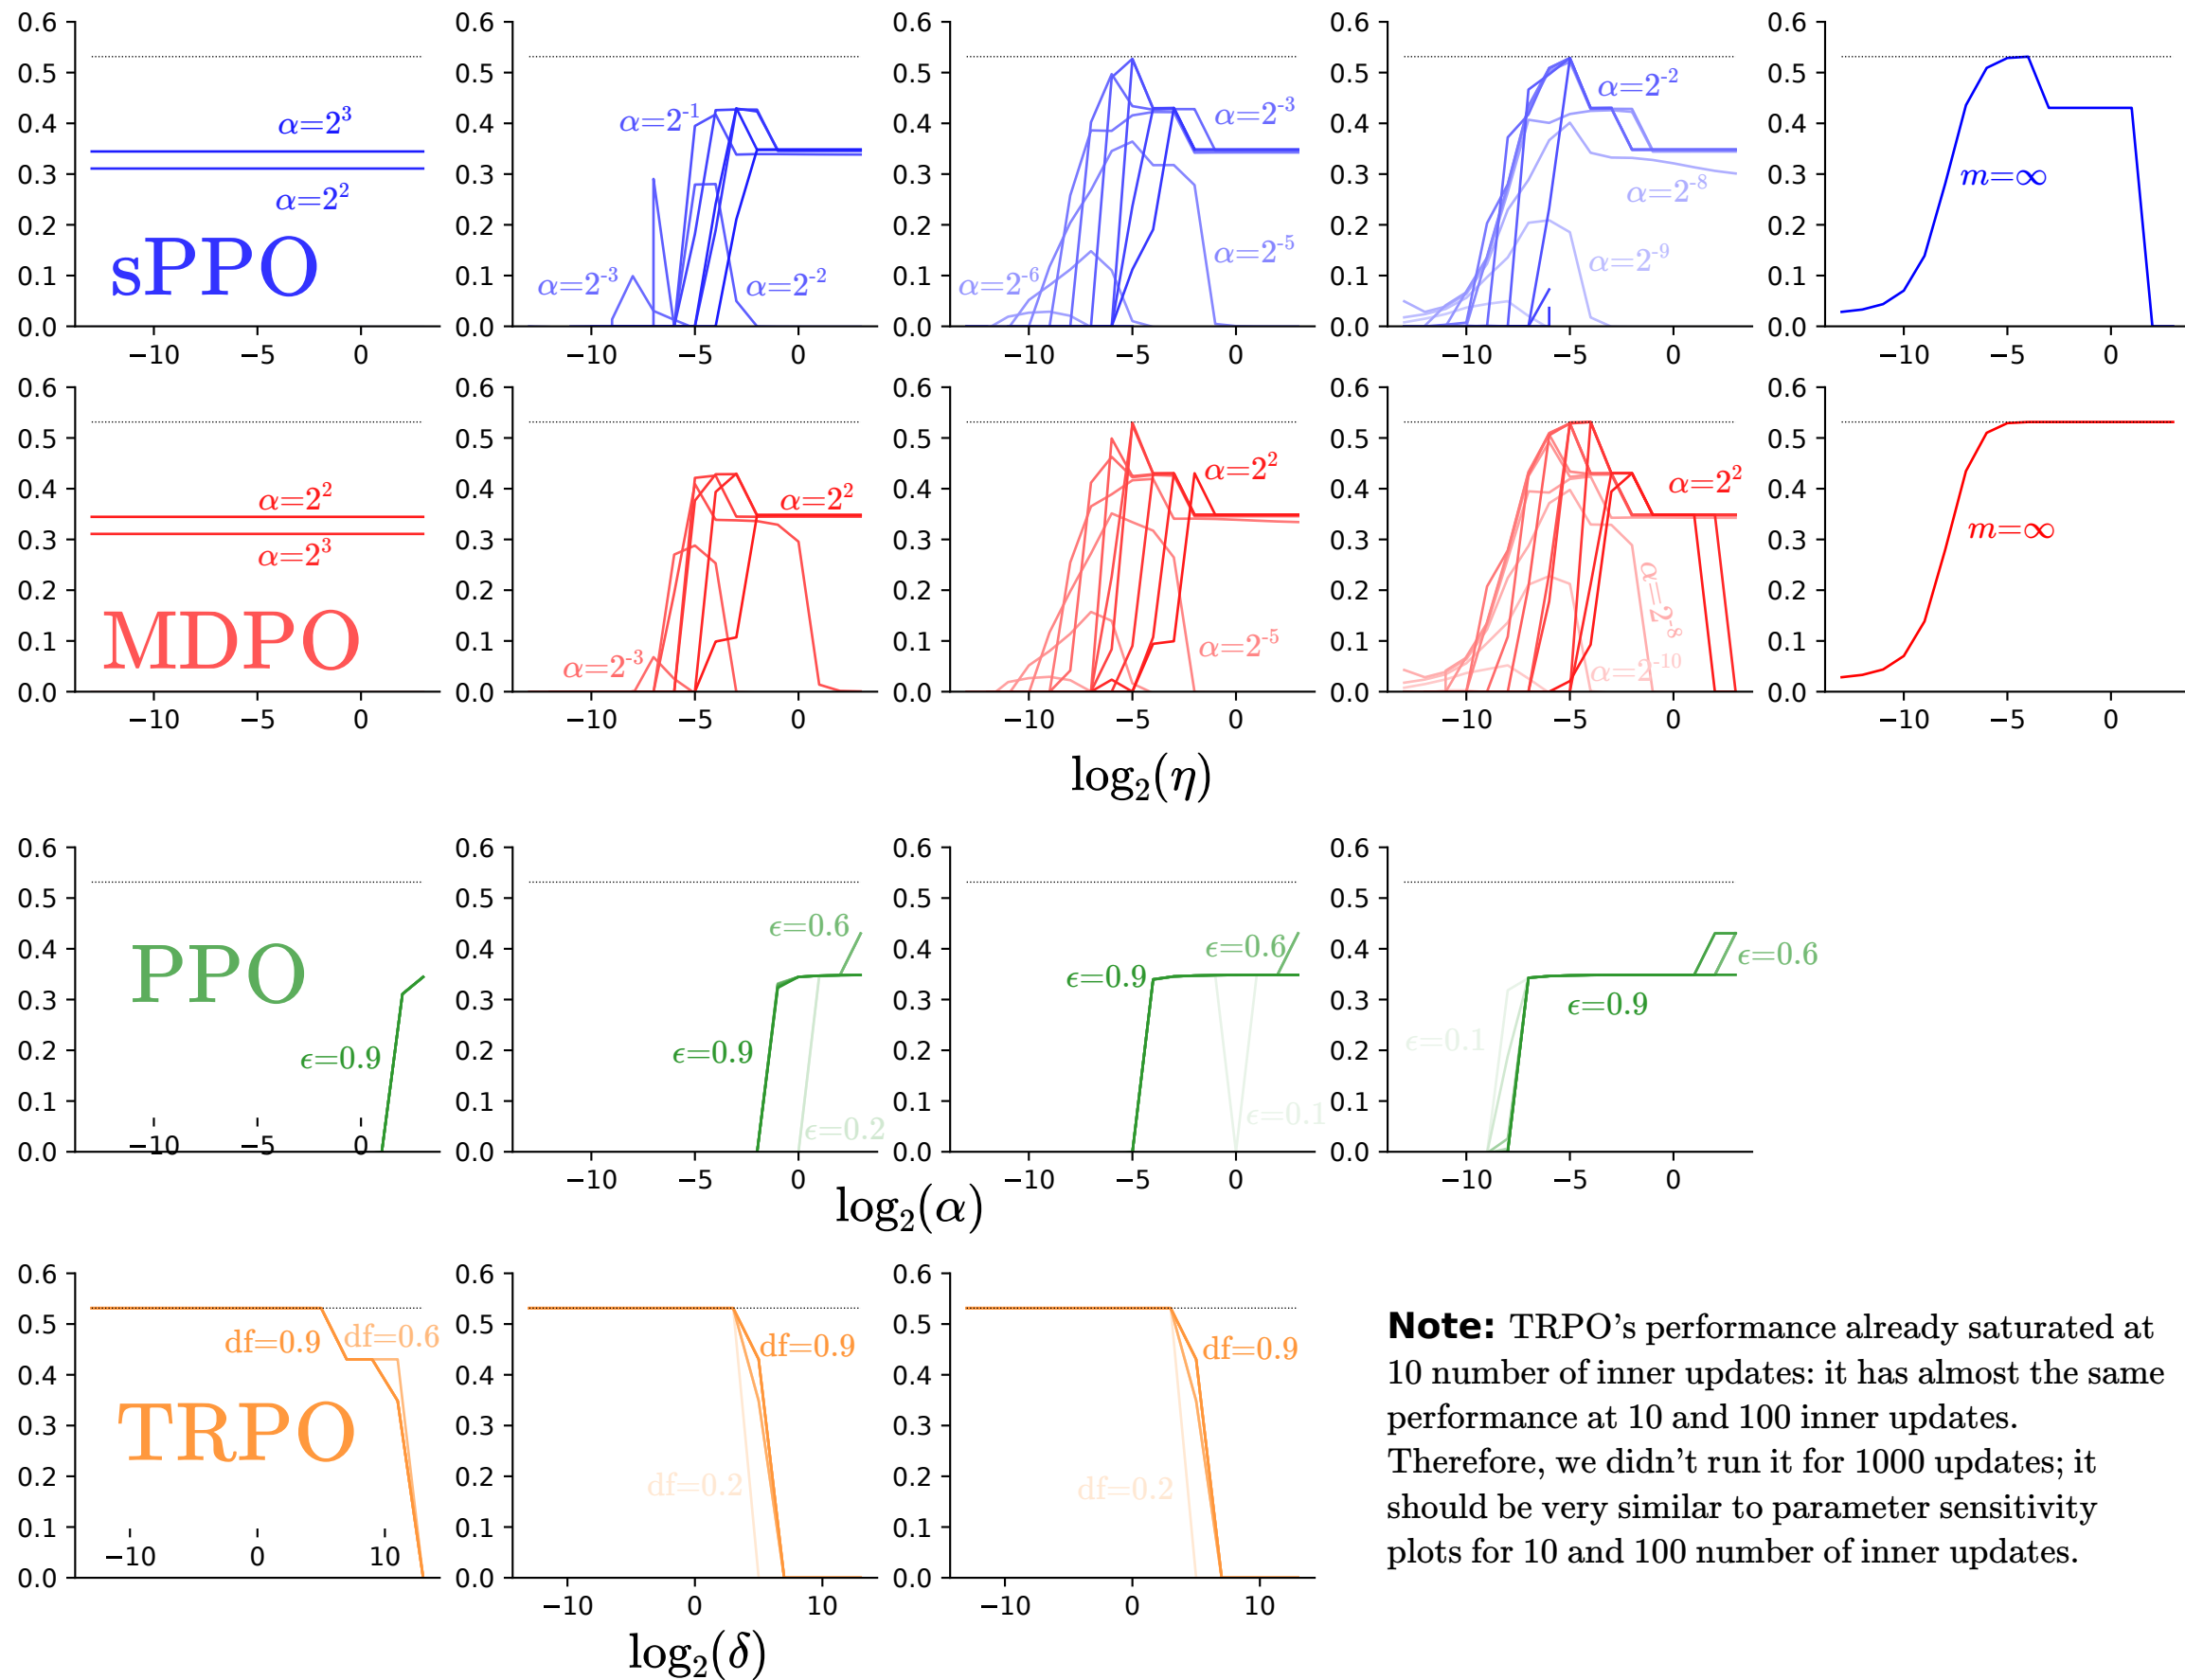


of inner updates: 1 10 100 1000 Analytical

Final Value of \mathcal{J} (PG objective)



Note: TRPO's performance already saturated at 10 number of inner updates: it has almost the same performance at 10 and 100 inner updates. Therefore, we didn't run it for 1000 updates; it should be very similar to parameter sensitivity plots for 10 and 100 number of inner updates.